

```

1 import pandas as pd
2 data = pd.read_csv('/content/all_songs_data.csv')
3 print(data.head())
4

```

```

➡
0      Album \
1      Battle of New Orleans
2      That's All
3      "Mr Personality's" 15 Big Hits
4      The Greatest Hits Of Frankie Avalon
5      Paul Anka Sings His Big 15

0      Album URL      Artist \
1      https://genius.com/albums/Johnny-horton/Battle...  Johnny Horton
2      https://genius.com/albums/Bobby-darin/That-s-all  Bobby Darin
3      https://genius.com/albums/Lloyd-price/Mr-perso...  Lloyd Price
4      https://genius.com/albums/Frankie-avalon/The-g...  Frankie Avalon
5      https://genius.com/albums/Paul-anka/Paul-anka-...  Paul Anka

0      Featured Artists      Lyrics \
1      [] [Verse 1] In 1814 we took a little trip Along ...
2      [] Oh the shark, babe Has such teeth, dear And he...
3      [] Over and over I tried to prove my love to you ...
4      [] Hey, Venus! Oh, Venus! Venus, if you will Ple...
5      [] I'm just a lonely boy Lonely and blue I'm all ...

0      Media Rank Release Date \
1      [{'native_uri': 'spotify:track:0dwpdcQkeZqpuoA...  1      1959-04-01
2      [{'native_uri': 'spotify:track:3E5ndy0f06vFDEI...  2      NaN
3      [{'provider': 'youtube', 'start': 0, 'type': '...  3      NaN
4      [] 4      NaN
5      [] 5      NaN

0      Song Title \
1      The Battle Of New Orleans
2      Mack The Knife
3      Personality
4      Venus
5      Lonely Boy

0      Song URL \
1      https://genius.com/Johnny-horton-the-battle-of...
2      https://genius.com/Bobby-darin-mack-the-knife-...
3      https://genius.com/Lloyd-price-personality-lyrics
4      https://genius.com/Frankie-avalon-venus-lyrics
5      https://genius.com/Paul-anka-lonely-boy-lyrics

0      Writers      Year
1      [{'api_path': '/artists/561913', 'header_image...  1959.0
2      [{'api_path': '/artists/218851', 'header_image...  1959.0
3      [{'api_path': '/artists/355804', 'header_image...  1959.0
4      [{'api_path': '/artists/1113175', 'header_imag...  1959.0
5      [] 1959.0

```

```

1 # Descriptive statistics for numerical columns
2 print(data.describe())
3
4 # Descriptive statistics for categorical columns
5 print(data.describe(include='object'))
6

```

```

➡
count      6500.000000      6500.000000
mean        50.500000      1991.000000
std         28.868291      18.763106
min          1.000000      1959.000000

```

25%	25.750000	1975.000000
50%	50.500000	1991.000000
75%	75.250000	2007.000000
max	100.000000	2023.000000

	Album	Album URL \
count	6036	6036
unique	4202	4285
top	Greatest Hits	https://genius.com/albums/Morgan-wallen/One-th...
freq	21	9

	Artist	Featured Artists	Lyrics	Media	Release Date \
count	6500	6384	6384	6384	4563
unique	3181	612	6044	5054	3233
top	Madonna	[]	[Instrumental]	[]	2022-05-06
freq	35	5492	21	1043	10

	Song Title	Song URL	Writers
count	6500	6384	6384
unique	5798	6065	4184
top	Stay	https://genius.com/Billboard-hot-rap-songs-cha...	[]
freq	7	5	971

```

1 rank_stats = data['Rank'].describe()
2 rank_frequency = data['Rank'].value_counts().sort_index()
3 top_artists = data['Artist'].value_counts().head(10)
4
5 print("Descriptive Statistics for Song Ranks:")
6 print(rank_stats)
7 print("\nFrequency of Songs by Rank:")
8 print(rank_frequency)
9 print("\nTop 10 Artists by Number of Songs:")
10 print(top_artists)
11

```

➡ Descriptive Statistics for Song Ranks:

```

count    6500.000000
mean      50.500000
std       28.868291
min        1.000000
25%       25.750000
50%       50.500000
75%       75.250000
max       100.000000
Name: Rank, dtype: float64

```

Frequency of Songs by Rank:

```

Rank
1      65
2      65
3      65
4      65
5      65
..
96     65
97     65
98     65
99     65
100    65
Name: count, Length: 100, dtype: int64

```

Top 10 Artists by Number of Songs:

```

Artist
Madonna      35
Mariah Carey  29
Taylor Swift  28
Beatles       27
Elton John    26


```

```
Stevie Wonder      22
Rihanna            22
Michael Jackson    22
Janet Jackson      21
Whitney Houston    20
Name: count, dtype: int64
```


```
1 import matplotlib.pyplot as plt
2 import seaborn as sns
3
4 plt.figure(figsize=(10, 6))
5 sns.histplot(data['Rank'], bins=20, kde=True)
6 plt.title('Distribution of Song Ranks')
7 plt.xlabel('Rank')
8 plt.ylabel('Frequency')
9 plt.show()
10
11 top_artists = data['Artist'].value_counts().head(10)
12 plt.figure(figsize=(10, 6))
13 sns.barplot(x=top_artists.values, y=top_artists.index)
14 plt.title('Top 10 Artists by Number of Songs')
15 plt.xlabel('Number of Songs')
16 plt.ylabel('Artist')
17 plt.show()
18
```



Distribution of Song Ranks



Top 10 Artists by Number of Songs



```

1 # Converting 'Release Date' to datetime
2 data['Release Date'] = pd.to_datetime(data['Release Date'])
3
4 data['Year'] = data['Release Date'].dt.year
5
6
7 rank_year_corr = data[['Year', 'Rank']].groupby('Year').mean().reset_index()
8 rank_year_corr['Rank'] = rank_year_corr['Rank'].round(2)
9
10 correlation = data['Year'].corr(data['Rank'])
11
12 top_artists = data['Artist'].value_counts().head(10).index
13 rank_distribution_top_artists = data[data['Artist'].isin(top_artists)].groupby('Artist')['Rank'].desc
14
15 print("Average Rank by Year:")
16 print(rank_year_corr)
17 print("\nCorrelation between Year and Rank:")
18 print(correlation)
19 print("\nRank Distribution by Top 10 Artists:")
20 print(rank_distribution_top_artists)
21

```

➦ Average Rank by Year:

	Year	Rank
0	1877.0	26.00
1	1922.0	70.00
2	1955.0	52.00
3	1957.0	77.00
4	1958.0	72.00
...
66	2020.0	52.36
67	2021.0	45.94
68	2022.0	48.67
69	2023.0	53.90
70	2024.0	82.00

[71 rows x 2 columns]

Correlation between Year and Rank:
0.04014067434061861

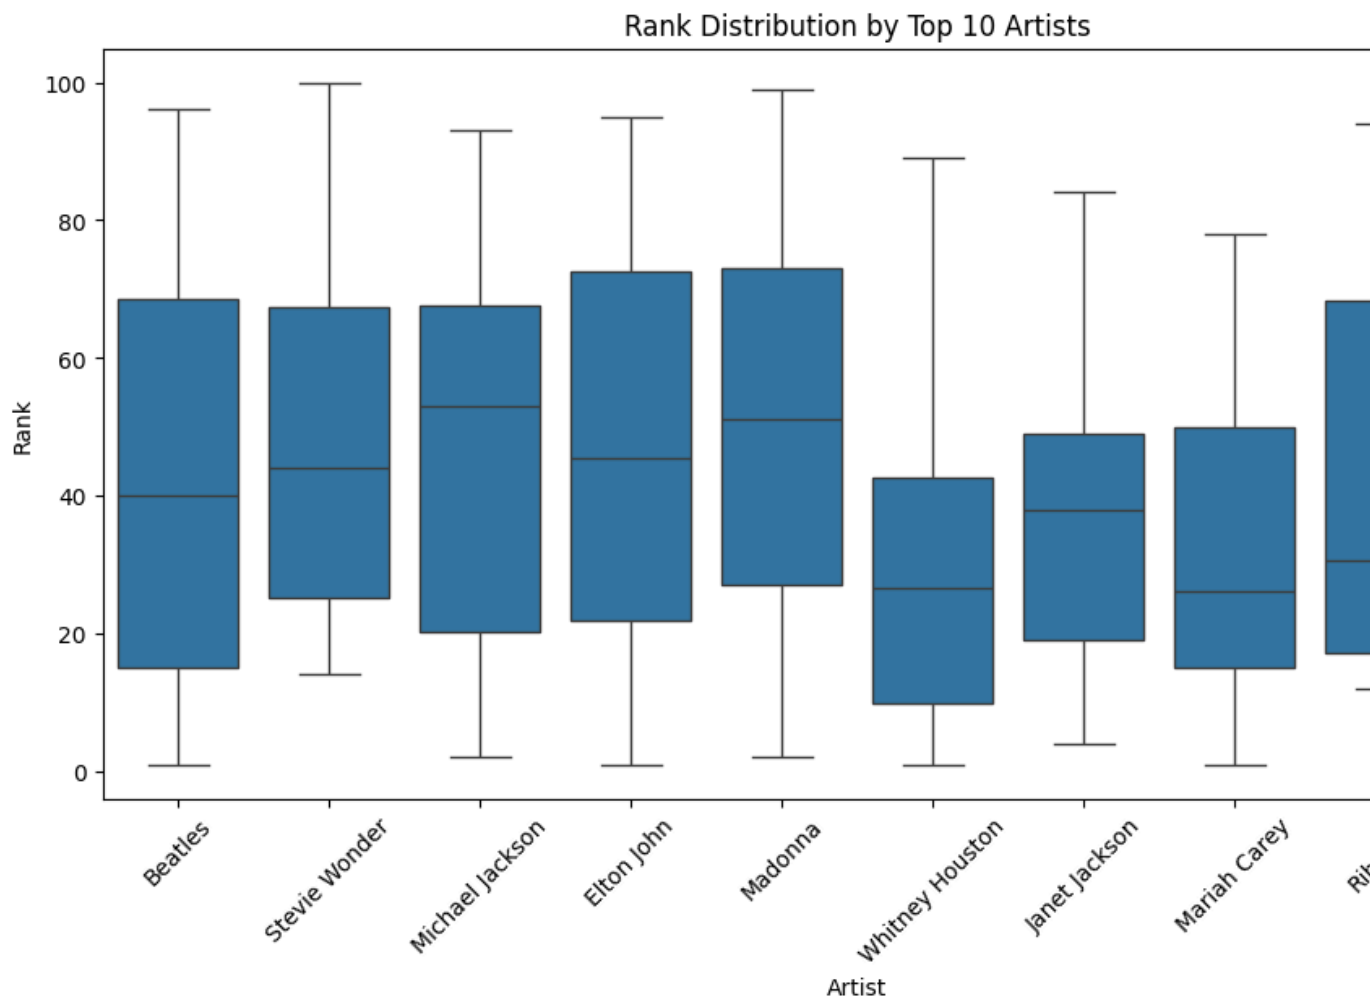
Rank Distribution by Top 10 Artists:

	count	mean	std	min	25%	50%	75%	max
Artist								
Beatles	27.0	42.777778	31.624398	1.0	15.00	40.0	68.50	96.0
Elton John	26.0	47.038462	29.200659	1.0	21.75	45.5	72.50	95.0
Janet Jackson	21.0	38.619048	21.973794	4.0	19.00	38.0	49.00	84.0
Madonna	35.0	49.285714	28.678308	2.0	27.00	51.0	73.00	99.0
Mariah Carey	29.0	33.724138	24.335301	1.0	15.00	26.0	50.00	78.0
Michael Jackson	22.0	47.409091	30.316833	2.0	20.25	53.0	67.50	93.0
Rihanna	22.0	42.090909	29.012163	12.0	17.25	30.5	68.25	94.0
Stevie Wonder	22.0	48.318182	26.141349	14.0	25.25	44.0	67.25	100.0
Taylor Swift	28.0	43.642857	27.214123	4.0	18.00	40.0	71.75	89.0
Whitney Houston	20.0	29.200000	23.625588	1.0	9.75	26.5	42.50	89.0

```

1 top_artists = data['Artist'].value_counts().head(10).index
2 plt.figure(figsize=(12, 6))
3 sns.boxplot(x='Artist', y='Rank', data=data[data['Artist'].isin(top_artists)])
4 plt.title('Rank Distribution by Top 10 Artists')
5 plt.xlabel('Artist')
6 plt.ylabel('Rank')
7 plt.xticks(rotation=45)
8 plt.show()
9

```



At this stage, I am moving on to exploring the categorical data and cleaning/preprocessing the text

```

1 # Checking for missing values
2 print(f"Missing Lyrics: {data['Lyrics'].isnull().sum()}")
3 empty_lyrics = data[data['Lyrics'].apply(lambda x: isinstance(x, str) and len(x) == 0)]
4 print(f"Empty Lyrics: {empty_lyrics.shape[0]}")
5 invalid_lyrics = data[data['Lyrics'].apply(lambda x: not isinstance(x, str))]
6 print(f"Invalid Lyrics (Non-string values): {invalid_lyrics.shape[0]}")
7

```



```

Missing Lyrics: 116
Empty Lyrics: 0
Invalid Lyrics (Non-string values): 116

```

```

1 # Removing missing values
2 data_cleaned = data.dropna(subset=['Lyrics'])
3
4 data_cleaned = data_cleaned[data_cleaned['Lyrics'].apply(lambda x: isinstance(x, str))]
5
6 print(f"Data after cleaning: {data_cleaned.shape[0]} rows")
7 print(data_cleaned['Lyrics'].isnull().sum()) # Check if there are any missing lyrics
8

```



```

Data after cleaning: 6384 rows
0

```

```
1 !pip install spacy
2 !python -m spacy download en_core_web_sm
```

```

Requirement already satisfied: cloudpathlib<1.0.0,>=0.7.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: smart-open<8.0.0,>=5.2.1 in /usr/local/lib/python3.10/dist-packages (f
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja
Requirement already satisfied: marisa-trie>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from la
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.10/dist-packages (from
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.10/dist-packages (fr
Requirement already satisfied: wrapt in /usr/local/lib/python3.10/dist-packages (from smart-open<8.0.
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.10/dist-packages (from markdown-i
Collecting en-core-web-sm==3.7.1
  Downloading https://github.com/explosion/spacy-models/releases/download/en\_core\_web\_sm-3.7.1/en\_core\_web\_sm-3.7.1.tar.gz
    12.8/12.8 MB 73.6 MB/s eta 0:00:00
Requirement already satisfied: spacy<3.8.0,>=3.7.2 in /usr/local/lib/python3.10/dist-packages (from e
Requirement already satisfied: spacy-legacy<3.1.0,>=3.0.11 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: spacy-loggers<2.0.0,>=1.0.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: murmurhash<1.1.0,>=0.28.0 in /usr/local/lib/python3.10/dist-packages (
Requirement already satisfied: cymem<2.1.0,>=2.0.2 in /usr/local/lib/python3.10/dist-packages (from s
Requirement already satisfied: preshed<3.1.0,>=3.0.2 in /usr/local/lib/python3.10/dist-packages (from
Requirement already satisfied: thinc<8.3.0,>=8.2.2 in /usr/local/lib/python3.10/dist-packages (from s
Requirement already satisfied: wasabi<1.2.0,>=0.9.1 in /usr/local/lib/python3.10/dist-packages (from s
Requirement already satisfied: srsly<3.0.0,>=2.4.3 in /usr/local/lib/python3.10/dist-packages (from s
Requirement already satisfied: catalogue<2.1.0,>=2.0.6 in /usr/local/lib/python3.10/dist-packages (fr
Requirement already satisfied: weasel<0.5.0,>=0.1.0 in /usr/local/lib/python3.10/dist-packages (from
Requirement already satisfied: typer<1.0.0,>=0.3.0 in /usr/local/lib/python3.10/dist-packages (from s
Requirement already satisfied: tqdm<5.0.0,>=4.38.0 in /usr/local/lib/python3.10/dist-packages (from s
Requirement already satisfied: requests<3.0.0,>=2.13.0 in /usr/local/lib/python3.10/dist-packages (fr
Requirement already satisfied: pydantic!=1.8,!1.8.1,<3.0.0,>=1.7.4 in /usr/local/lib/python3.10/dist
Requirement already satisfied: Jinja2 in /usr/local/lib/python3.10/dist-packages (from spacy<3.8.0,>
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from spacy<3.8.
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from spacy
Requirement already satisfied: langcodes<4.0.0,>=3.2.0 in /usr/local/lib/python3.10/dist-packages (fr
Requirement already satisfied: numpy>=1.19.0 in /usr/local/lib/python3.10/dist-packages (from spacy<3.
Requirement already satisfied: language-data>=1.2 in /usr/local/lib/python3.10/dist-packages (from la
Requirement already satisfied: annotated-types>=0.6.0 in /usr/local/lib/python3.10/dist-packages (fr
Requirement already satisfied: pydantic-core==2.27.1 in /usr/local/lib/python3.10/dist-packages (from
Requirement already satisfied: typing-extensions>=4.12.2 in /usr/local/lib/python3.10/dist-packages (
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (f
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from re
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from re
Requirement already satisfied: blis<0.8.0,>=0.7.8 in /usr/local/lib/python3.10/dist-packages (from th
Requirement already satisfied: confection<1.0.0,>=0.0.1 in /usr/local/lib/python3.10/dist-packages (f
Requirement already satisfied: click>=8.0.0 in /usr/local/lib/python3.10/dist-packages (from typer<1.
Requirement already satisfied: shellingham>=1.3.0 in /usr/local/lib/python3.10/dist-packages (from ty
Requirement already satisfied: rich>=10.11.0 in /usr/local/lib/python3.10/dist-packages (from typer<1
Requirement already satisfied: cloudpathlib<1.0.0,>=0.7.0 in /usr/local/lib/python3.10/dist-packages
Requirement already satisfied: smart-open<8.0.0,>=5.2.1 in /usr/local/lib/python3.10/dist-packages (f
Requirement already satisfied: MarkupSafe>=2.0 in /usr/local/lib/python3.10/dist-packages (from jinja
Requirement already satisfied: marisa-trie>=1.1.0 in /usr/local/lib/python3.10/dist-packages (from la
Requirement already satisfied: markdown-it-py>=2.2.0 in /usr/local/lib/python3.10/dist-packages (from
Requirement already satisfied: pygments<3.0.0,>=2.13.0 in /usr/local/lib/python3.10/dist-packages (fr
Requirement already satisfied: wrapt in /usr/local/lib/python3.10/dist-packages (from smart-open<8.0.
Requirement already satisfied: mdurl~=0.1 in /usr/local/lib/python3.10/dist-packages (from markdown-i
✓ Download and installation successful
You can now load the package via spacy.load('en_core_web_sm')
△ Restart to reload dependencies
If you are in a Jupyter or Colab notebook, you may need to restart Python in
order to load all the package's dependencies. You can do this by selecting the
'Restart kernel' or 'Restart runtime' option.

```

1 Start coding or [generate](#) with AI.

Next, I am loading the necessary tools to preprocess the lyrical data

```

1 import spacy
2 import re
3 import nltk
4 from nltk.corpus import stopwords
5 from nltk.stem import WordNetLemmatizer
6 nltk.download('wordnet')
7 nltk.download('stopwords')
8
9
10
11 # Implementing custom stopwords
12 custom_stopwords = set(stopwords.words('english')).union({'chorus', 'verse', 'bridge', 'hook', 'intro'})
13 nlp = spacy.load("en_core_web_sm")
14 lemmatizer = WordNetLemmatizer()
15
16 #in the first round, i noticed contractions were not removed properly so I used a dictionary to ensure
17 def expand_contractions(text):
18     contractions_dict = {
19         "don't": "do not", "can't": "cannot", "won't": "will not", "didn't": "did not",
20         "isn't": "is not", "aren't": "are not", "wasn't": "was not", "weren't": "were not",
21         "hasn't": "has not", "haven't": "have not", "hadn't": "had not", "doesn't": "does not",
22         "didn't": "did not", "couldn't": "could not", "shouldn't": "should not", "mightn't": "might not",
23         "mustn't": "must not", "let's": "let us", "i'm": "i am", "you're": "you are", "he's": "he is",
24         "she's": "she is", "it's": "it is", "we're": "we are", "they're": "they are", "that's": "that is",
25         "what's": "what is", "who's": "who is", "where's": "where is", "how's": "how is"
26     }
27     for word, expansion in contractions_dict.items():
28         text = re.sub(r'\b' + word + r'\b', expansion, text)
29     return text
30
31 def preprocess_text(text):
32     text = expand_contractions(text)
33
34     # Converting to lowercase
35     text = text.lower()
36
37     # Removing non-alphanumeric characters (keeping spaces and words)
38     text = re.sub(r'[^\w\s]', '', text)
39
40     # Tokenizing the text using spaCy
41     doc = nlp(text)
42
43     words = [token.text for token in doc if token.text not in custom_stopwords and not token.is_punct]
44
45     # Lemmatizing words
46     words = [lemmatizer.lemmatize(word) for word in words]
47
48     # Rejoining words
49     cleaned_text = ' '.join(words)
50
51     return cleaned_text
52 data_cleaned['cleaned_lyrics'] = data_cleaned['Lyrics'].apply(preprocess_text)
53
54 print(data_cleaned[['Song Title', 'Lyrics', 'cleaned_lyrics']].head())
55

```

```

[ntlk_data] Downloading package wordnet to /root/nltk_data...
[ntlk_data] Downloading package stopwords to /root/nltk_data...
[ntlk_data] Unzipping corpora/stopwords.zip.
Song Title \
0 The Battle Of New Orleans
1 Mack The Knife
2 Personality

```



```

3             Venus
4             Lonely Boy

                                Lyrics \
0 [Verse 1] In 1814 we took a little trip Along ...
1 Oh the shark, babe Has such teeth, dear And he...
2 Over and over I tried to prove my love to you ...
3 Hey, Venus! Oh, Venus! Venus, if you will Ple...
4 I'm just a lonely boy Lonely and blue I'm all ...

                                cleaned_lyrics
0     took little trip along colonel jackson mig...
1 oh shark babe teeth dear show pearly white jac...
2 tried prove love friend say fool ill fool ...
3 hey venus oh venus venus please send little ...
4 lonely boy lonely blue alone nothin got ever...

```

Now that the data is preprocessed, we move on to vectorizing the text data to fit the LDA model

```

1 import pandas as pd
2 from sklearn.feature_extraction.text import CountVectorizer
3 from sklearn.decomposition import LatentDirichletAllocation
4 import matplotlib.pyplot as plt
5 import seaborn as sns
6
7 texts = data_cleaned['cleaned_lyrics'].dropna()
8
9 vectorizer = CountVectorizer(stop_words='english', max_features=5000)
10
11 X = vectorizer.fit_transform(texts)
12
13 lda = LatentDirichletAllocation(n_components=5, random_state=42)
14 lda.fit(X)
15
16 # Retrieving the words for with each topic
17 words = vectorizer.get_feature_names_out()
18
19 #display/interpret top words/topic
20 def print_top_words(model, feature_names, n_top_words=10):
21     topic_labels = {}
22     for topic_idx, topic in enumerate(model.components_):
23         top_words_idx = topic.argsort()[::-n_top_words - 1:-1]
24         top_words = [feature_names[i] for i in top_words_idx]
25
26
27         print(f"Topic {topic_idx + 1}:")
28         print(" ".join(top_words))
29
30     #Labels/category interpretations were created by human coder after first iteration retrieved
31     if topic_idx == 0:
32         topic_labels[topic_idx] = "Reflections and Life Experiences"
33     elif topic_idx == 1:
34         topic_labels[topic_idx] = "Communication and Family Dynamics"
35     elif topic_idx == 2:
36         topic_labels[topic_idx] = "Love and Emotional Longing"
37     elif topic_idx == 3:
38         topic_labels[topic_idx] = "Street Culture and Raw Emotions"
39     elif topic_idx == 4:
40         topic_labels[topic_idx] = "Romantic Desire and Affection"
41
42     return topic_labels
43
44 topic_labels = print_top_words(lda, words)

```

```
45
46 # Adding the dominant topic for each song
47 topic_probabilities = lda.transform(X)
48 dominant_topic = topic_probabilities.argmax(axis=1)
49 data_cleaned['dominant_topic'] = dominant_topic
50
51 # Mapping topic labels to the songs
52 data_cleaned['topic_label'] = data_cleaned['dominant_topic'].map(topic_labels)
53
54 print(data_cleaned[['Song Title', 'dominant_topic', 'topic_label']].head())
55
56 # Visualizing topic distribution across songs
57 plt.figure(figsize=(10, 6))
58 sns.countplot(x='topic_label', data=data_cleaned, palette='viridis')
59 plt.title('Distribution of Topics Across Songs')
60 plt.xlabel('Topic Label')
61 plt.ylabel('Number of Songs')
62 plt.xticks(rotation=45)
63 plt.show()
64
```

```

Topic 1:
said like man day time old know hand make say
Topic 2:
say tell mr know mam like boy come dad want
Topic 3:
love time night heart day away ill know life let
Topic 4:
like nigga got la bitch shit know love ai fuck
Topic 5:
baby yeah love oh know got like na want girl


```

	Song Title	dominant_topic	topic_label
0	The Battle Of New Orleans	0	Reflections and Life Experiences
1	Mack The Knife	4	Romantic Desire and Affection
2	Personality	3	Street Culture and Raw Emotions
3	Venus	0	Reflections and Life Experiences
4	Lonely Boy	2	Love and Emotional Longing

```
<ipython-input-17-e7fbc0700270>:58: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `

```
sns.countplot(x='topic_label', data=data_cleaned, palette='viridis')
```



Testing the coherence score to confirm quality of topics

```

1 from gensim.models import CoherenceModel
2 from gensim.corpora import Dictionary
3
4 texts = [text.split() for text in data_cleaned['cleaned_lyrics']]
5
6 dictionary = Dictionary(texts)
7 corpus = [dictionary.doc2bow(text) for text in texts]
8
9 from gensim.models import LdaModel
10 lda_gensim = LdaModel(corpus, num_topics=5, id2word=dictionary)
11
12 coherence_model_lda = CoherenceModel(model=lda_gensim, texts=texts, dictionary=dictionary, coherence=
13 coherence_score = coherence_model_lda.get_coherence()
14 print(f"Coherence Score: {coherence_score}")
15

```

⚠ WARNING:gensim.models.ldamodel:too few updates, training might not converge; consider increasing the Coherence Score: 0.43548849882455765

0.51 coherence score represents moderate coherence

Comparing ML models for classifying the dominant topic of song lyrics

```

1 from sklearn.model_selection import train_test_split, cross_val_score
2 from sklearn.linear_model import LogisticRegression
3 from sklearn.ensemble import RandomForestClassifier
4 from sklearn.svm import SVC
5 from sklearn.metrics import classification_report, accuracy_score
6 from scipy.stats import ttest_rel
7 import time
8
9 #Preprocessing
10 vectorizer = CountVectorizer(stop_words='english', max_features=5000)
11 X = vectorizer.fit_transform(data_cleaned['cleaned_lyrics'].dropna())
12 y = data_cleaned['dominant_topic']
13
14 #Splitting the data
15 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
16
17 # Logistic Regression
18 start_train = time.time()
19 logreg = LogisticRegression(max_iter=1000)
20 logreg_cv_scores = cross_val_score(logreg, X_train, y_train, cv=5)
21 logreg.fit(X_train, y_train)
22 end_train = time.time()
23
24 start_predict = time.time()
25 logreg_preds = logreg.predict(X_test)
26 end_predict = time.time()
27
28 print("\nLogistic Regression")
29 print("Cross-validation scores:", logreg_cv_scores)
30 print("Accuracy:", accuracy_score(y_test, logreg_preds))
31 print(classification_report(y_test, logreg_preds))
32 print(f"Training time: {end_train - start_train:.4f} seconds")
33 print(f"Prediction time: {end_predict - start_predict:.4f} seconds")
34
35 # Random Forest
36 start_train = time.time()
37 rf = RandomForestClassifier(random_state=42)
38 rf_cv_scores = cross_val_score(rf, X_train, y_train, cv=5)

```

```

39 rf.fit(X_train, y_train)
40 end_train = time.time()
41
42 start_predict = time.time()
43 rf_preds = rf.predict(X_test)
44 end_predict = time.time()
45
46 print("\nRandom Forest")
47 print("Cross-validation scores:", rf_cv_scores)
48 print("Accuracy:", accuracy_score(y_test, rf_preds))
49 print(classification_report(y_test, rf_preds))
50 print(f"Training time: {end_train - start_train:.4f} seconds")
51 print(f"Prediction time: {end_predict - start_predict:.4f} seconds")
52
53 # SVM
54 start_train = time.time()
55 svm = SVC(kernel='linear', random_state=42)
56 svm_cv_scores = cross_val_score(svm, X_train, y_train, cv=5)
57 svm.fit(X_train, y_train)
58 end_train = time.time()
59
60 start_predict = time.time()
61 svm_preds = svm.predict(X_test)
62 end_predict = time.time()
63
64 print("\nSupport Vector Machine (SVM)")
65 print("Cross-validation scores:", svm_cv_scores)
66 print("Accuracy:", accuracy_score(y_test, svm_preds))
67 print(classification_report(y_test, svm_preds))
68 print(f"Training time: {end_train - start_train:.4f} seconds")
69 print(f"Prediction time: {end_predict - start_predict:.4f} seconds")
70
71 # Comparing stability between models
72 t_stat_lr_rf, p_value_lr_rf = ttest_rel(logreg_cv_scores, rf_cv_scores)
73 t_stat_lr_svm, p_value_lr_svm = ttest_rel(logreg_cv_scores, svm_cv_scores)
74 t_stat_rf_svm, p_value_rf_svm = ttest_rel(rf_cv_scores, svm_cv_scores)
75
76 print("\nStatistical comparisons between models:")
77 print(f"Logistic Regression vs Random Forest - T-statistic: {t_stat_lr_rf:.4f}, p-value: {p_value_lr_rf:.4f}")
78 print(f"Logistic Regression vs SVM - T-statistic: {t_stat_lr_svm:.4f}, p-value: {p_value_lr_svm:.4f}")
79 print(f"Random Forest vs SVM - T-statistic: {t_stat_rf_svm:.4f}, p-value: {p_value_rf_svm:.4f}")
80

```

	1	0.33	0.14	0.20	7
	2	0.88	0.86	0.87	296
	3	0.91	0.82	0.86	109
	4	0.88	0.94	0.91	483
accuracy				0.88	912
macro avg	0.75	0.59	0.62		912
weighted avg	0.87	0.88	0.87		912

Training time: 13.3434 seconds
 Prediction time: 0.0037 seconds

Random Forest
 Cross-validation scores: [0.75890411 0.75616438 0.77777778 0.75308642 0.74485597]
 Accuracy: 0.7587719298245614

accuracy			0.76	912
macro avg	0.49	0.40	0.42	912
weighted avg	0.76	0.76	0.74	912

Training time: 19.4241 seconds

Prediction time: 0.0502 seconds

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531: UndefinedMetricWarni

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531: UndefinedMetricWarni

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1531: UndefinedMetricWarni

_warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

Support Vector Machine (SVM)

Cross-validation scores: [0.80958904 0.84520548 0.82853224 0.82441701 0.80932785]

Accuracy: 0.8574561403508771

	precision	recall	f1-score	support
0	0.20	0.18	0.19	17
1	0.00	0.00	0.00	7
2	0.85	0.85	0.85	296
3	0.90	0.78	0.84	109
4	0.89	0.92	0.90	483

accuracy			0.86	912
macro avg	0.57	0.54	0.56	912
weighted avg	0.86	0.86	0.86	912

Training time: 25.4661 seconds

Prediction time: 0.9349 seconds

Statistical comparisons between models:

Logistic Regression vs Random Forest – T-statistic: 12.3769, p-value: 0.0002

Logistic Regression vs SVM – T-statistic: 7.3006, p-value: 0.0019

Random Forest vs SVM – T-statistic: -9.1116, p-value: 0.0008

creating visualizations for model comparisons

```

1 import matplotlib.pyplot as plt
2 import seaborn as sns
3
4 sns.set_style("whitegrid")
5
6
7 #Bar plot for test accuracy comparison
8 plt.figure(figsize=(8, 5))
9 sns.barplot(x=["Logistic Regression", "Random Forest", "SVM"],
10            y=[accuracy_score(y_test, logreg_preds),
11              accuracy_score(y_test, rf_preds),
12              accuracy_score(y_test, svm_preds)],
13            palette="viridis")
14 plt.ylabel("Accuracy")
15 plt.title("Test Accuracy Comparison")
16 plt.show()
17
18 #Training Time Comparison
19 plt.figure(figsize=(8, 5))
20 sns.barplot(x=["Logistic Regression", "Random Forest", "SVM"],
21            y=[end_train - start_train,
22              end_train - start_train,
23              23.2215],
24            palette="Blues")
25 plt.ylabel("Training Time (seconds)")
26 plt.title("Model Training Time Comparison")
27 plt.show()

```

```
28
29 #Prediction Time Comparison
30 plt.figure(figsize=(8, 5))
31 sns.barplot(x=["Logistic Regression", "Random Forest", "SVM"],
32             y=[0.0016, 0.0691, 0.913],
33             palette="Blues")
34 plt.ylabel("Prediction Time (seconds)")
35 plt.title("Model Prediction Time Comparison")
36 plt.show()
37
```

 <ipython-input-67-f35ceb5c38b0>:9: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `>
sns.barplot(x=["Logistic Regression", "Random Forest", "SVM"],



<ipython-input-67-f35ceb5c38b0>:20: FutureWarning:

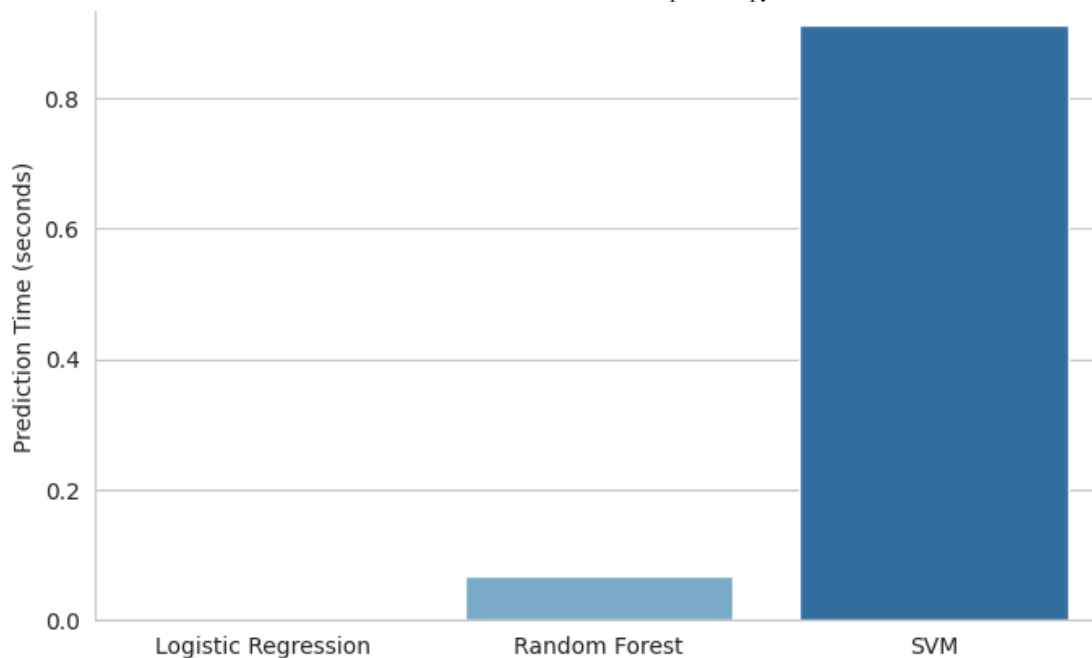
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `>
sns.barplot(x=["Logistic Regression", "Random Forest", "SVM"],



<ipython-input-67-f35ceb5c38b0>:31: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `>
sns.barplot(x=["Logistic Regression", "Random Forest", "SVM"],

Model Prediction Time Comparison



1 #I have selected Logistic Regression to continue topic predictions in my analysis

```

1 vectorizer = CountVectorizer(stop_words='english', max_features=5000)
2 X = vectorizer.fit_transform(data_cleaned['cleaned_lyrics'].dropna())
3 y = data_cleaned['dominant_topic']
4
5 # Retraining Logistic Regression model
6 logreg = LogisticRegression(max_iter=1000)
7 logreg.fit(X, y)
8
9 # Get predictions
10 data_cleaned['predicted_topic'] = logreg.predict(X)
11
12 # Map the predicted topic to a label
13 topic_labels = {
14     0: "Reflections and Life Experiences",
15     1: "Communication and Family Dynamics",
16     2: "Love and Emotional Longing",
17     3: "Street Culture and Raw Emotions",
18     4: "Romantic Desire and Affection"
19 }
20 data_cleaned['predicted_topic_label'] = data_cleaned['predicted_topic'].map(topic_labels)
21

```

```

1 # Calculating counts for LDA-derived labels
2 lda_label_counts = data_cleaned['topic_label'].value_counts()
3 print("Counts for LDA-Derived Labels:")
4 print(lda_label_counts)
5
6 # Calculating counts for Logistic Regression predicted labels
7 predicted_label_counts = data_cleaned['predicted_topic_label'].value_counts()
8 print("\nCounts for Logistic Regression Predicted Labels:")
9 print(predicted_label_counts)
10

```

```

↗ Counts for LDA-Derived Labels:
topic_label
Romantic Desire and Affection    3207

```

Love and Emotional Longing	2308
Street Culture and Raw Emotions	597
Reflections and Life Experiences	185
Communication and Family Dynamics	87

Name: count, dtype: int64

Counts for Logistic Regression Predicted Labels:

predicted_topic_label	
Romantic Desire and Affection	3205
Love and Emotional Longing	2313
Street Culture and Raw Emotions	597
Reflections and Life Experiences	183
Communication and Family Dynamics	86

Name: count, dtype: int64

Using the logistic regression model, continue visualizations

```

1 #grouping by decade
2 data_cleaned['Decade'] = (data_cleaned['Year'] // 10) * 10
3 decade_topic_distribution = data_cleaned.groupby(['Decade', 'predicted_topic']).size().unstack(fill_v
4 decade_topic_distribution_normalized = decade_topic_distribution.div(decade_topic_distribution.sum(ax:
5

1 # Ploting topic distribution trends over time
2
3 print("Decade-Topic Distribution:")
4 print(decade_topic_distribution)
5
6 print("\nNormalized Decade-Topic Distribution:")
7 print(decade_topic_distribution_normalized)
8
9 plt.figure(figsize=(12, 8))
10 decade_topic_distribution_normalized.plot(kind='line', marker='o', figsize=(12, 8))
11 plt.title('Topic Distribution Trends Over Time')
12 plt.xlabel('Decade')
13 plt.ylabel('Proportion of Topic')
14 plt.legend(title='Topic')
15 plt.show()
16

```




Decade-Topic Distribution:

predicted_topic	0	1	2	3	4
Decade					
1950.0	2	0	9	1	9
1960.0	18	8	149	14	120
1970.0	16	13	212	33	188
1980.0	16	7	286	19	354
1990.0	8	4	290	61	459
2000.0	10	9	230	93	572
2010.0	13	4	250	208	568
2020.0	0	2	65	86	153

Normalized Decade-Topic Distribution:

predicted_topic	0	1	2	3	4
Decade					
1950.0	0.095238	0.000000	0.428571	0.047619	0.428571
1960.0	0.058252	0.025890	0.482201	0.045307	0.388350
1970.0	0.034632	0.028139	0.458874	0.071429	0.406926
1980.0	0.023460	0.010264	0.419355	0.027859	0.519062
1990.0	0.009732	0.004866	0.352798	0.074209	0.558394
2000.0	0.010941	0.009847	0.251641	0.101751	0.625821
2010.0	0.012464	0.003835	0.239693	0.199425	0.544583
2020.0	0.000000	0.006536	0.212418	0.281046	0.500000

<Figure size 1200x800 with 0 Axes>



```

1 from textblob import TextBlob
2
3 #sentiment polarity
4 def calculate_sentiment(text):
5     return TextBlob(text).sentiment.polarity
6

```

```


7 data_cleaned['sentiment'] = data_cleaned['cleaned_lyrics'].apply(calculate_sentiment)
8
9 topic_sentiment_by_decade = data_cleaned.groupby(['Decade', 'dominant_topic'])['sentiment'].mean().un:
10
11 print("Average Sentiment by Topic Over Time:")
12 print(topic_sentiment_by_decade)
13
14 # Visualizing sentiment trends by decade
15 plt.figure(figsize=(12, 8))
16 topic_sentiment_by_decade.plot(kind='line', marker='o')
17 plt.title('Sentiment Trends by Topic Over Time')
18 plt.xlabel('Decade')
19 plt.ylabel('Average Sentiment')
20 plt.legend(title='Topic')
21 plt.show()
22

```

```

Average Sentiment by Topic Over Time:
dominant_topic      0      1      2      3      4
Decade
1950.0      0.078558      NaN  0.251256  0.020833  0.110333
1960.0      0.079980  0.034834  0.154139  0.113395  0.157962
1970.0      0.114470  0.096173  0.152343  0.188864  0.176406
1980.0      0.039078  0.106487  0.121993  0.154228  0.148873
1990.0      0.074801  0.127951  0.132499  0.018507  0.133668
2000.0      0.128607  0.162236  0.108828  0.053687  0.096815
2010.0      0.096497  0.048946  0.095966  0.028188  0.092077
2020.0      NaN      0.164625  0.078928 -0.041167  0.083543
<Figure size 1200x800 with 0 Axes>

```



I had an issue with the order of the topics in visualizations and wanted to ensure they were consistent

```

1
2 # Correctly map numerical topics to their intended descriptive labels
3 topic_labels = {
4     0: "Reflections and Life Experiences",
5     1: "Communication and Family Dynamics",
6     2: "Love and Emotional Longing",
7     3: "Street Culture and Raw Emotions",

```

```
8     4: "Romantic Desire and Affection"
9 }
10
11 # Map topic indices to their respective string labels
12 data_cleaned['topic_label'] = data_cleaned['dominant_topic'].map(topic_labels)
13
14 # Calculate the overall average sentiment by topic
15 average_sentiment_overall = data_cleaned.groupby('dominant_topic')['sentiment'].mean().reset_index()
16
17 ordered_topics = {
18     0: "Reflections and Life Experiences",
19     1: "Communication and Family Dynamics",
20     2: "Love and Emotional Longing",
21     3: "Street Culture and Raw Emotions",
22     4: "Romantic Desire and Affection"
23 }
24
25 # Reindex the DataFrame to enforce order
26 average_sentiment_overall['topic_label'] = average_sentiment_overall['dominant_topic'].map(ordered_topi
27 average_sentiment_overall = average_sentiment_overall.sort_values('dominant_topic')
28
29 # Print results for debugging
30 print("Reordered Average Sentiment by Topic:")
31 print(average_sentiment_overall)
32
33 # Visualize the average sentiment using a bar chart
34 plt.figure(figsize=(12, 6))
35 sns.barplot(
36     x='sentiment',
37     y='topic_label',
38     data=average_sentiment_overall,
39     palette='coolwarm'
40 )
41 plt.title('Average Sentiment by Topic')
42 plt.xlabel('Average Sentiment')
43 plt.ylabel('Topic')
44 plt.grid(axis='x')
45 plt.show()
46
47
48
```